

# Assignment Number: 11

Name: Divya Shah; Branch: I.T (T.E.); Roll Number: 115

21/09/2023

**Aim:** To study and configure Firewalls using IP tables.

**LO mapped:** LO6

**Theory:**

## Firewall-

A firewall is a system designed to prevent unauthorized access to or from a private network. You can implement a firewall in either hardware or software form, or a combination of both. Generally, the firewall has two network interfaces: one for the external side of the network, one for the internal side. Its purpose is to control what traffic is allowed to traverse from one side to the other. As the most basic level, firewalls can block traffic intended for particular IP addresses or server ports.

TCP network traffic moves around a network in packets, which are containers that consist of a packet header—this contains control information such as source and destination addresses, and packet sequence information—and the data (also known as a payload). While the control information in each packet helps to ensure that its associated data gets delivered properly, the elements it contains also provides firewalls a variety of ways to match packets against firewall rules.

## Types Of Firewalls:

There are three basic types of firewalls which are mentioned below:

- I. **Packet filtering**, or stateless, firewalls work by inspecting individual packets in isolation. As such, they are unaware of connection state and can only allow or deny packets based on individual packet headers.
- II. **Stateful firewalls** are able to determine the connection state of packets, which makes them much more flexible than stateless firewalls. They work by collecting related packets until the connection state can be determined before any firewall rules are applied to the traffic.
- III. **Stateful firewalls** are able to determine the connection state of packets, which makes them much more flexible than stateless firewalls. They work by collecting related packets until the connection state can be determined before any firewall rules are applied to the traffic.

## Basics Of IPTABLES:

Iptables is a firewall, installed by default on all official Ubuntu distributions (Ubuntu, Kubuntu, Xubuntu). When you install Ubuntu, iptables is there, but it allows all traffic by default.

The rules in Iptables are written to deal 3 different scenarios:

1. Those packets entering your machine that are destined for your machine. (INPUT)
2. Those packets leaving your machine. (OUTPUT)
3. Those packets entering your machine, but are destined for another machine and will pass through your machine (FORWARD).

In Iptables, these scenarios are referred to as INPUT, OUTPUT, and FORWARD, respectively.

Once the traffic type has been specified, three actions may be taken:

1. ACCEPT allows packets to pass through the firewall.
2. DROP ignores the packet and sends no response to the request.
3. REJECT ignores the packet, but responds to the request with a packet denied message.

### Basic Commands-

Command - `sudo iptables -L` (`-L` – List the current filter rules.)

```

vyashah@vyashah:~$ sudo iptables -L
(sudo) password for dvyashah:
Chain INPUT (policy ACCEPT)
target     prot opt source                destination

Chain FORWARD (policy DROP)
target     prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination

Chain DOCKER-USER (1 references)
target     prot opt source                destination
DOCKER-USER all -- anywhere             anywhere

Chain DOCKER-ISOLATION-STAGE-1 (1 references)
target     prot opt source                destination
ACCEPT     all -- anywhere             anywhere             ctstate RELATED,ESTABLISHED
DOCKER     all -- anywhere             anywhere
ACCEPT     all -- anywhere             anywhere
ACCEPT     all -- anywhere             anywhere

Chain DOCKER-ISOLATION-STAGE-2 (1 references)
target     prot opt source                destination
DROP       all -- anywhere             anywhere
RETURN     all -- anywhere             anywhere

Chain DOCKER (1 references)
target     prot opt source                destination
RETURN     all -- anywhere             anywhere
  
```

As you can see, we have our three default chains (INPUT, OUTPUT, and FORWARD). We also can see each chain's default policy (each chain has ACCEPT as its default policy). We also see some column headers, but we don't see any actual rules. This is because Ubuntu doesn't ship with a default rule set.

### Basic Iptables Options:

- -A - Append this rule to a rule chain. Valid chains for what we're doing are INPUT, FORWARD and OUTPUT, but we mostly deal with INPUT in this tutorial, which affects only incoming traffic.
- -p - The connection protocol used.

- `--dport` - The destination port(s) required for this rule. A single port may be given, or a range may be given as start: end, which will match all ports from start to end, inclusive.
- `-j` - Jump to the specified target. By default, iptables allows four targets:
- `ACCEPT` - Accept the packet and stop processing rules in this chain.
- `REJECT` - Reject the packet and notify the sender that we did so, and stop processing rules in this chain.
- `DROP` - Silently ignore the packet, and stop processing rules in this chain.
- `LOG` - Log the packet, and continue processing more rules in this chain. Allows the use of the `--log-prefix` and `--log-level` options.
- `-i` - Only match if the packet is coming in on the specified interface.
- `-I` - Inserts a rule. Takes two options, the chain to insert the rule into, and the rule number it should be.
- `-I INPUT 5` would insert the rule into the INPUT chain and make it the 5<sup>th</sup> rule in the list.
- `-v` - Display more information in the output. Useful for if you have rules that look similar without using `-v`.
- `-s --source` - address[/mask] source specification
- `-d --destination` - address[/mask] destination specification
- `-o --out-interface` - output name[+] network interface name ([+] for wildcard)

#### Allowing Incoming Traffic on Specific Ports -

You could start by blocking traffic, but you might be working over SSH, where you would need to allow SSH before blocking everything else. To allow incoming traffic on the default SSH port (22), you could tell iptables to allow all TCP traffic on that port to come in.

```
sudo iptables -A INPUT -p tcp --dport ssh -j ACCEPT
```

Referring back to the list above, you can see that this tells iptables:

- append this rule to the input chain (`-A INPUT`) so we look at incoming traffic
- check to see if it is TCP (`-p tcp`).
- if so, check to see if the input goes to the SSH port (`--dport ssh`).
- if so, accept the input (`-j ACCEPT`).

Now, let's allow all incoming web traffic:

```
sudo iptables -A INPUT -p tcp --dport 80 -j ACCEPT
```

*We have specifically allowed tcp traffic to the ssh and web ports, but as we have not blocked anything, all traffic can still come in.*

#### Blocking Traffic

Once a decision is made to accept a packet, no more rules affect it. As our rules allowing ssh and web traffic come first, as long as our rule to block all traffic comes after them, we can still accept the traffic we want. All we need to do is put the rule to block all traffic at the end.

```
sudo iptables -A INPUT -j DROP
sudo iptables -L
```

Because we didn't specify an interface or a protocol, any traffic for any port on any interface is blocked, except for web and SSH.

```
divyashah@DivyaShah:~$ sudo iptables -A INPUT -j DROP
divyashah@DivyaShah:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source                destination            tcp dpt:ssh
ACCEPT     tcp  --  anywhere              anywhere               tcp dpt:ssh
DROP       all  --  anywhere              anywhere

Chain FORWARD (policy DROP)
target     prot opt source                destination
DOCKER-USER all  --  anywhere              anywhere
DOCKER-ISOLATION-STAGE-1 all  --  anywhere              anywhere
ACCEPT     all  --  anywhere              anywhere               ctstate RELATED,ESTABLISHED
DOCKER     all  --  anywhere              anywhere
ACCEPT     all  --  anywhere              anywhere
ACCEPT     all  --  anywhere              anywhere

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination

Chain DOCKER (1 references)
target     prot opt source                destination

Chain DOCKER-ISOLATION-STAGE-1 (1 references)
target     prot opt source                destination
DOCKER-ISOLATION-STAGE-2 all  --  anywhere              anywhere
RETURN     all  --  anywhere              anywhere

Chain DOCKER-ISOLATION-STAGE-2 (1 references)
target     prot opt source                destination
DROP       all  --  anywhere              anywhere
RETURN     all  --  anywhere              anywhere

Chain DOCKER-USER (1 references)
target     prot opt source                destination
RETURN     all  --  anywhere              anywhere
divyashah@DivyaShah:~$
```

### Editing iptables

The only problem with our setup so far is that even the loopback port is blocked. We could have written the drop rule for just eth0 by specifying -i eth0, but we could also add a rule for the loopback. If we append this rule, it will come too late - after all the traffic has been dropped. We need to insert this rule before that. Since this is a lot of traffic, we'll insert it as the first rule so it's processed first.

```
sudo iptables -I INPUT 1 -i lo -j
ACCEPT sudo iptables -L
```

```

Chain DOCKER-ISOLATION-STAGE-2 (1 references)
target prot opt source destination
DROP all -- anywhere anywhere
RETURN all -- anywhere anywhere

Chain DOCKER-USER (1 references)
target prot opt source destination
RETURN all -- anywhere anywhere
vyyashah@vyyashah:~$ sudo iptables -I INPUT 1 -i lo -j ACCEPT
vyyashah@vyyashah:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target prot opt source destination
ACCEPT all -- anywhere anywhere
ACCEPT tcp -- anywhere anywhere tcp dpt:ssh
DROP all -- anywhere anywhere
DROP all -- anywhere anywhere

Chain FORWARD (policy DROP)
target prot opt source destination
DOCKER-USER all -- anywhere anywhere
DOCKER-ISOLATION-STAGE-1 all -- anywhere anywhere
ACCEPT all -- anywhere anywhere ctstate RELATED,ESTABLISHED
ACCEPT all -- anywhere anywhere
ACCEPT all -- anywhere anywhere
ACCEPT all -- anywhere anywhere

Chain OUTPUT (policy ACCEPT)
target prot opt source destination
Chain DOCKER (1 references)
target prot opt source destination
Chain DOCKER-ISOLATION-STAGE-1 (1 references)
target prot opt source destination
DOCKER-ISOLATION-STAGE-2 all -- anywhere anywhere
RETURN all -- anywhere anywhere
Chain DOCKER-ISOLATION-STAGE-2 (1 references)
target prot opt source destination
DROP all -- anywhere anywhere
RETURN all -- anywhere anywhere
Chain DOCKER-USER (1 references)
target prot opt source destination
RETURN all -- anywhere anywhere
vyyashah@vyyashah:~$

```

We will list iptables in greater detail.

*sudo iptables -L -v*

You can now see a lot more information. This rule is actually very important, since many programs use the loopback interface to communicate with each other.

```

vyyashah@vyyashah:~$ sudo iptables -L -v
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
pkts bytes target prot opt in out source destination
14 1173 ACCEPT all -- lo any anywhere anywhere
0 0 ACCEPT tcp -- any any anywhere anywhere tcp dpt:ssh
200 35438 DROP all -- any any anywhere anywhere
0 0 DROP all -- any any anywhere anywhere

Chain FORWARD (policy DROP 0 packets, 0 bytes)
pkts bytes target prot opt in out source destination
0 0 DOCKER-USER all -- any any anywhere anywhere
0 0 DOCKER-ISOLATION-STAGE-1 all -- any any anywhere anywhere
0 0 DOCKER all -- any docker0 anywhere anywhere ctstate RELATED,ESTABLISHED
0 0 DOCKER all -- docker0 docker0 anywhere anywhere
0 0 ACCEPT all -- docker0 !docker0 anywhere anywhere
0 0 ACCEPT all -- docker0 docker0 anywhere anywhere

Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
pkts bytes target prot opt in out source destination

Chain DOCKER (1 references)
pkts bytes target prot opt in out source destination

Chain DOCKER-ISOLATION-STAGE-1 (1 references)
pkts bytes target prot opt in out source destination
0 0 DOCKER-ISOLATION-STAGE-2 all -- docker0 !docker0 anywhere anywhere
0 0 RETURN all -- any any anywhere anywhere

Chain DOCKER-ISOLATION-STAGE-2 (1 references)
pkts bytes target prot opt in out source destination
0 0 DROP all -- any docker0 anywhere anywhere
0 0 RETURN all -- any any anywhere anywhere

Chain DOCKER-USER (1 references)
pkts bytes target prot opt in out source destination
0 0 RETURN all -- any any anywhere anywhere
vyyashah@vyyashah:~$

```

### Allow traffic on ICMP port

*sudo iptables -A INPUT -p icmp -j ACCEPT*

Now we have to list rules again...

*sudo iptables -L*

```
Chain DOCKER-ISOLATION-STAGE-2 (1 references)
target prot opt source destination
0 0 DROP all -- any docker anywhere
0 0 RETURN all -- any anywhere

Chain DOCKER-USER (1 references)
target prot opt source destination
0 0 RETURN all -- any anywhere
divyashah@DivyaShah:~$ sudo iptables -A INPUT -p icmp -j ACCEPT
divyashah@DivyaShah:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target prot opt source destination
ACCEPT all -- anywhere anywhere
ACCEPT tcp -- anywhere anywhere tcp dpt:ssh
DROP all -- anywhere anywhere
ACCEPT icmp -- anywhere anywhere

Chain FORWARD (policy DROP)
target prot opt source destination
DOCKER-USER all -- anywhere anywhere
DOCKER-ISOLATION-STAGE-1 all -- anywhere anywhere
ACCEPT all -- anywhere anywhere ctstate RELATED,ESTABLISHED
ACCEPT all -- anywhere anywhere
ACCEPT all -- anywhere anywhere

Chain OUTPUT (policy ACCEPT)
target prot opt source destination

Chain DOCKER (1 references)
target prot opt source destination

Chain DOCKER-ISOLATION-STAGE-1 (1 references)
target prot opt source destination
RETURN all -- anywhere anywhere

Chain DOCKER-ISOLATION-STAGE-2 (1 references)
target prot opt source destination
DROP all -- anywhere anywhere
RETURN all -- anywhere anywhere

Chain DOCKER-USER (1 references)
target prot opt source destination
RETURN all -- anywhere anywhere
divyashah@DivyaShah:~$
```

clearing all rules

```
sudo iptables -F
```

```
sudo iptables -L
```

```
divyashah@DivyaShah:~$ sudo iptables -F
divyashah@DivyaShah:~$ sudo iptables -L
Chain INPUT (policy ACCEPT)
target prot opt source destination

Chain FORWARD (policy DROP)
target prot opt source destination

Chain OUTPUT (policy ACCEPT)
target prot opt source destination

Chain DOCKER (0 references)
target prot opt source destination

Chain DOCKER-ISOLATION-STAGE-1 (0 references)
target prot opt source destination

Chain DOCKER-ISOLATION-STAGE-2 (0 references)
target prot opt source destination

Chain DOCKER-USER (0 references)
target prot opt source destination
divyashah@DivyaShah:~$
```

### Dropping icmp packets

Try to ping your neighbour machine

```
ping 192.168.92.17
```

You can see the response packets received. Now block incoming icmp packets from the neighbour using command:

```
sudo iptables -A INPUT -p icmp -j DROP
```

List the rule:

```
sudo iptables -L
```

Try to ping your neighbour machine again

```
ping 192.168.92.17
```

You cannot see receive ICMP echo reply packets...

Now try to restrict outgoing icmp packets by adding rule

```
sudo iptables -A OUTPUT -p icmp -j DROP List the rule:
```

```
sudo iptables -L
```

Now try to ping neighbour,

```
ping 192.168.92.17
```

Flush all rules and try to ping neighbour, Blocking TCP port traffic will not allow u to browse the Internet

```
sudo iptables -F INPUT -p tcp -j DROP List the rule:
```

```
sudo iptables -L
```

Now try to access the internet. You can't. Flush the rule n then try to access internet you can.

**Blocking ICMP packets from specific source machine:**

```
sudo iptables -A INPUT -s 192.168.92.17 -p icmp -j DROP
```

```
sudo iptables -L
```

```
ping 192.168.92.17
```

– – –does not allow (192.168.92.17 can not send u icmp packets)

ping any other machine:

```
ping 192.168.92.11
```

– – –allowed (192.168.92.11 can send u icmp packets)

**Types of iptables:**

## I. IPTABLES TABLES and CHAINS

IPTables has the following 4 built-in tables.

### 1. Filter Table

Filter is default table for iptables. So, if you don't define you own table, you'll be using filter table. Iptables's filter table has the following built-in chains.

- INPUT chain – Incoming to firewall. For packets coming to the local server.
- OUTPUT chain – Outgoing from firewall. For packets generated locally and going out of the local server.

- FORWARD chain – Packet for another NIC on the local server. For packets routed through the local server.

Type the following command and see the

result

```
sudo iptables -t filter -L
```

## 2. NAT table

Iptable's NAT table has the following built-in chains.

- PREROUTING chain – Alters packets before routing. i.e Packet translation happens immediately after the packet comes to the system (and before routing). This helps to translate the destination ip address of the packets to something that matches the routing on the local server. This is used for DNAT (destination NAT).
- POSTROUTING chain – Alters packets after routing. i.e Packet translation happens when the packets are leaving the system. This helps to translate the source ip address of the packets to something that might match the routing on the destination server.

This is used for SNAT (source NAT).

- OUTPUT chain – NAT for locally generated packets on the firewall.

Type the following command and see the result

```
sudo iptables -t nat -L
```

## 3. Mangle table

Iptables's Mangle table is for specialized packet alteration. This alters QOS bits in the TCP header. Mangle table has the following built-in chains.

- PREROUTING chain
- OUTPUT chain
- FORWARD chain
- INPUT chain
- POSTROUTING chain

Type the following command and see the result

```
sudo iptables -t nat -L
```

## 4. Raw table

Iptable's Raw table is for configuration exemptions. Raw table has the following built-in chains.



- PREROUTING chain
- OUTPUT chain

**Conclusion:** We implemented a studied Firewalls using IP tables.